

Source? I Made it Up: How Hallucinations Can Be Studied to Reveal Inherent Political Bias in LLMs

Pipeline combining generation, detection, and classification

- LLM outputs detect hallucinations classify bias

Use of multiple open-source LLMs for comparison

- Llama, Mistral, DeepSeek evaluated

Hallucinations vary significantly by model

- Frequency differs across models

Hallucinations occur equally across political inputs

- Similar rates for left vs right

Hallucinated content shows systematic left-leaning bias

- Skews left, especially on key topics

Currently pending conference acceptance

